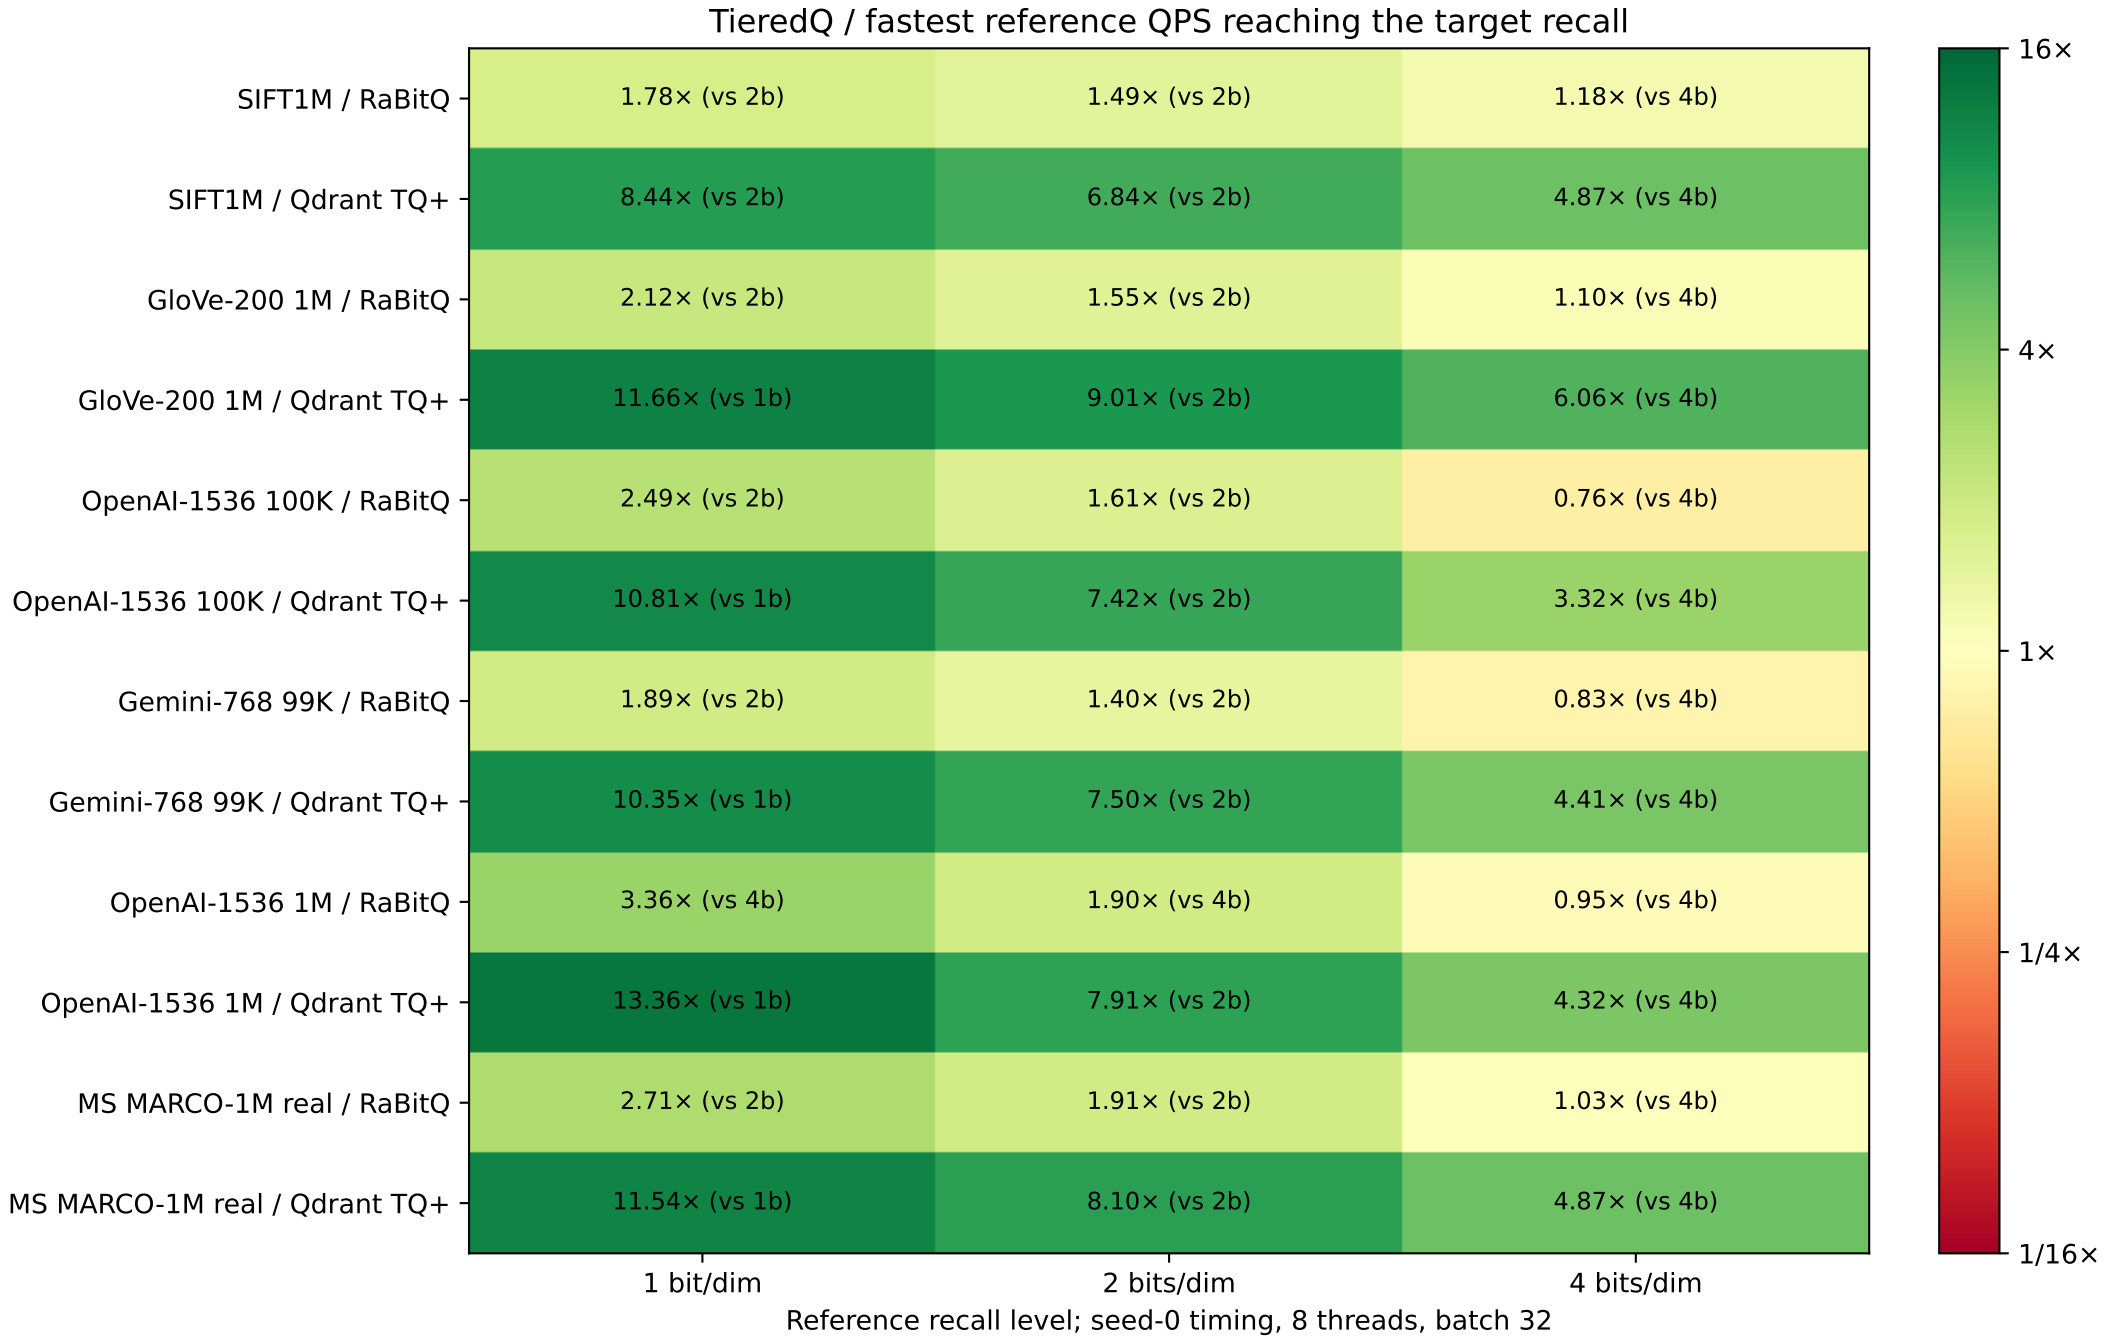
TieredQ / fastest reference QPS reaching the target recall
SIFT1M / RaBitQ
1.78× (vs 2b)
1.49× (vs 2b)
1.18× (vs 4b)
SIFT1M / Qdrant TQ+
8.44× (vs 2b)
6.84× (vs 2b)
4.87× (vs 4b)
GloVe-200 1M / RaBitQ
2.12× (vs 2b)
1.55× (vs 2b)
1.10× (vs 4b)
GloVe-200 1M / Qdrant TQ+
11.66× (vs 1b)
9.01× (vs 2b)
6.06× (vs 4b)
OpenAI-1536 100K / RaBitQ
2.49× (vs 2b)
1.61× (vs 2b)
0.76× (vs 4b)
OpenAI-1536 100K / Qdrant TQ+
10.81× (vs 1b)
7.42× (vs 2b)
3.32× (vs 4b)
Gemini-768 99K / RaBitQ
1.89× (vs 2b)
1.40× (vs 2b)
0.83× (vs 4b)
Gemini-768 99K / Qdrant TQ+
10.35× (vs 1b)
7.50× (vs 2b)
4.41× (vs 4b)
OpenAI-1536 1M / RaBitQ
3.36× (vs 4b)
1.90× (vs 4b)
0.95× (vs 4b)
OpenAI-1536 1M / Qdrant TQ+
13.36× (vs 1b)
7.91× (vs 2b)
4.32× (vs 4b)
MS MARCO-1M real / RaBitQ
2.71× (vs 2b)
1.91× (vs 2b)
1.03× (vs 4b)
MS MARCO-1M real / Qdrant TQ+
11.54× (vs 1b)
8.10× (vs 2b)
4.87× (vs 4b)
1 bit/dim
2 bits/dim
4 bits/dim
Reference recall level; seed-0 timing, 8 threads, batch 32
16×
4×
1×
1/4×
1/16×